Time	Given Name	Family Name	Type of Proposal	Topic or Name of proposed talk, workshop or project	Abstract of talk or project
Saturday March 14					
9:00 AM	Brian	Proffitt	User Meetup/Workshop - 3 hours	oVirt Workshop	We would like to host a full- or half-day oVirt user/admin meetup, colocated with the FOSSAsia event. We would provide speakers to fill the available time, and would just need a room with standard AV facilities.
12:00 PM					
1:00 PM	Leap	Sok	Talk of 25 minutes	IT Automation with Ansible	Software configuration management tools are gaining ground everywhere. We've all probably at least heard of puppet, chef or salt by now, however there's a new rising star: Ansible. In this talk we'll talk about the way Ansible approaches configuration management, software deployment, remote execution and other general IT tasks.
1:30 PM	Gerald	Goh	Talk of 25 minutes	Docker + Redhat + Openstack and the glue to hold them together	At Apvera we are focused on delivering services that make use of open source technologies to enhance how end users consume business applications. We primarily use Redhat + Docker + Openstack to enable enterprises to help consume application provisioning with simplicity and portability.
2:00 PM	Masahiro	Nakagawa	Talk of 25 minutes	Fluentd: Unified Logging Infrastructure	I plan to talk about my main open source project Fluentd. Fluentd is an open source data collector to unify logging infrastructure. Used at companies like Nintendo and Slideshare and integrated into projects like Kubernetes (Google's Docker orchestration framework), it is quickly becoming a popular tool to collect, process and manage log streams. In this talk, I will highlight the key features and design principles of Fluentd and explain why many programmers and ops engineers find it useful. Time permitting, I will share the project's roadmap and new features. In 2014, Fluentd maintainers have given talks at RedDotRubyConf, RubyConf, MongoDBWorld, OSCON and RubyKaigi. This year, we are expanding our reach to open source communities outside of Ruby (Fluentd is primarily written in Ruby).
2:30 PM	Sameer	Kumar	Talk of 25 minutes	10 Features your Developers are missing when 'Stuck' with proprietary database!	Database is core of any enterprise project. Be it the solution which a bank relies on for daily transactions or be the software which helps your mobile operator track your usage and send you a bill. Every application needs a data store and that is why every developer sooner or later has to be exposed to databases. I have seen mostly when I talk to developers about a database or relational database there are a few or rather just a name that would pop-up in their head. Proprietary databases have become so much a synonym to the term "RDBMS" that developer today hardly know about any other database. In this talk I will discuss about 10 such great features which your developers are missing if they are still stuck with Proprietary Database. I will be discussing 10 awesome features of PostgreSQL- Worlds most Advanced Open Source Database.
3:00 PM	Choon Ming	Goh	Talk of 25 minutes	Deploying OpenNebula, the Open Source Solution for Enterprise Cloud Solution	As enterprises expand their datacenters or internal infrastructure, it is very common for one to use commercial software such as VMWare or OSS solution such as Xen, KVM or OpenStack. However, the reality is that enterprises chose commercial products in the end because of the features made available to them. An open source cloud/infrastructure solution such as OpenNebula provides the features that can rivals VMWare whilst integrating with various other open source technologies will balance up the scale. The talk will show you how OpenNebula works and what are the features available and how it integrates with other Open Source solution in order to rival VMWare or other commercial products.
3:30 PM	Shivji Kumar	Jha	Talk of 25 minutes	Enhanced High Availability u	MySQL Replication provides a solution for High Availability and Re
4:00 PM	Ryusuke	Kajiyama	Talk of 25 minutes	MySQL Cluster as Transactional NoSQL	There's a lot of excitement around NoSQL Data Stores with the promise of simple access patterns, flexible schemas, scalability and High Availability. The downside comes in the form of losing ACID transactions, consistency, flexible queries and data integrity checks. What if you could have the best of both worlds? This session shows how MySQL Cluster provides simultaneous SQL and native NoSQL access to your data – whether a simple key-value API (Memcached), REST, JavaScript, Java or C++. You will hear how the MySQL Cluster architecture delivers in-memory real-time performance, 99.999% availability, on-line maintenance and linear, horizontal scalability through transparent auto-sharding.
4:30 PM	End				
Sunday March 15					
10:00 AM	Priyank	Shah	Talk of 25 minutes	Adopting Agile with Workshop	Adopting Agile with Workshop. Please find initial presentation hosted at http://www.slideshare.net/priyank17/adopting-agile-via-continuous-improvement-with-workshop 1. Agile Overview. 2. Scrum Overview. 3. Applicability/need/Usage of scrum. 4. Each artifacts of Scrum & all meeting. 5. Real challenges of Scrum 6. Estimation technique (Poker, Animal, T-Shirt Estimations) 7. Measurement using burn down chart. 8. Tools In market. 9. Workshop on sample project 10. face real situation and challenges during the workshop. 11. Q & A
10:30 AM	Amita	Sharma	Talk of 25 minutes	Testing - Best Practices in Open Souce World	1. How to contribute in improving quality of open source software 2. How to create test harness for open source software. 3. How to alter your software development model to produce bug-free code. 4. Altering your development model to match SW quality standards. 5. Study the technology, know the feature well and use your experience 6. Know the point when it is time to report a bug 7. How to troubleshoot that buggy situation.

Time	Given Name	Family Name	Type of Proposal	Topic or Name of proposed talk, workshop or project	Abstract of talk or project
11:00 AM	Piyush	Verma	Talk of 25 minutes	analysing your CDN usage	CDNlysis syncs Amazon Cloudfront logs, which are stored as compressed files in S3 bucket and streams them to InfluxDB, a time series Database with a powerful Querying API. CDNlysis is designed to work on millions and millions of records and as-it-happens updates of the CDN usage. You use this for: Understanding how the Bandwidth is being used. Finding out the most popular and most downloadable content. Generate trends for your most popular Videos, Audios, Slides etc. Understand geographical behaviour of the Requests. Amount of Bytes transferred to & fro the Cloudfront distributions. Find out the most profitable referrer from where your content is being accessed. etc.
11:30 AM	Satoshi	Nagayasu	Talk of 25 minutes	PostgreSQL 9.4 and Beyond	Recently, PostgreSQL has succeeded in getting popularity among web developers. One of the reasons is that PostgreSQL is growing not only as a traditional RDBMS, but also as an application development platform with fully supporting ACID properties. In this talk, I would like to introduce new features in the latest release 9.4 (JSON, Analytics, Logical Decoding, etc.) to take advantages of the latest PostgreSQL advancement, and share a quick look in the upcoming release 9.5.
12:00 PM					
1:00 PM	Maor	Lipchuk	Talk of 25 minutes	Disaster Recovery in oVirt	Even the best system administrator cannot always avoids any and every disaster that may plague his data center, but he should have a contingency plan to recover from one - and an administrator that manages his virtual data centers with oVirt is of course no different. This session will cover the new features introduced in oVirt 3.5.0 to handle such scenarios and will showcase how stringing together a set of building blocks can produce a well rounded solution for disaster scenarios.
1:30 PM	Ricky	Setyawan	Workshop of 1 hour	How to setup web scalable architecture with MySQL Fabric	MySQL fabric is a tool that was launched as part of MySQL Utilities 1.4. Fabric is written in python and includes a special library that implements all the functionalities provided. It seeks to make horizontal scale-out through sharding more accessible to users with growing data management requirements. This integrated framework supports management of large farms of MySQL servers, and includes support for sharding and high-availability. In this session, we will share how to setup a scalable web architecture with fabric.
2:00 PM	Brett	Gray	Talk of 25 minutes	MCollective, practical and useful	This presentation will provide details on MCollective including, architecture, security, customising, and writing your own MCollective Agents. We will also discuss and demonstrate how to use MCollective in a scaled environment.
2:30 PM	Mayank	Prasad	Talk of 25 minutes	MySQL Performance Schema - Monitoring MySQL Server with Perfromance Schema and enhancement in MySQL 5.7	The MySQL Performance Schema is a feature for monitoring Oracle's MySQL Server execution at a low level. This versatile and tightly integrated component collects performance and session data from various subsystems within the server during runtime with minimal impact on overall server performance. This feature was introduced in MySQL 5.5 and MySQL 5.6 was released with great enhancements in it for performance monitoring and tuning. In 5.7 milestone releases as well, there are multiple new functionalities have been added. This talk aims to give an overview of Performance Schema and how/what statistical information user is provided with using Performance Schema and how that information can be used for monitoring and trouble shooting to track an issue back to the relevant file and line of code in source file. This talk also explains how Performance Schema gives consolidated view of different server activities behind the scene with the multiple summary tables and how this can be tuned. This talk also aims to give a glimpse of mew features/monitoring information added in MySQL 5.7 Milestone releases. This talk also covers the manual configuration of Performance Schema to suite to individual requirement in a specific scenario.
3:00 PM	Chris	Kruppa	Talk of 25 minutes	How a lean and agile mindset can lead to more innovative products	Most innovative projects fail either because of wrong and not verified assumptions or because of fearing the risk to invest into insecure projects. The lean and agile mindset deals with those fears and offer certain practices to both develop an innovative product and to find fast the market fit. Therefore investments can be reduced and the chance for success increases by factors. In this workshop we will learn some basics of the lean/agile mindset and apply some of the practices introduced by it's leaders. The workshop will be very interactive: It is planned, that the participants will choose the practices they want to learn during this workshop.
3:30 PM	Lyle	Kozloff	Talk of 25 minutes	Configuration Management Systems on the Desktop	Tools like Puppet, Ansible and Chef are well known for their power in managing thousands of servers. They're not as well known for managing workstations. The unique nature of workstations makes this use case challenging in occasionally surprising ways. Lyle Kozloff has been managing hundreds of desktops and laptops in an education environment using Puppet Open Source for the past two years and will share some insights, provide some tips and show some code that will help anyone wanting to leverage a configuration management system in their own environment.
4:00 PM	Aditya	Patawari	Talk of 25 minutes	Running Project Atomic and Docker on Fedora	This session will focus on getting started with Docker with an introduction to Project Atomic. We'll discuss why Docker can be a better choice than LXC and virtual machines in many scenarios. We'll discuss rpm-ostree and its advantages followed by running a Fedora Atomic host feeding it cloud-init data. We'll also create custom docker images from Fedora image which can be shipped anywhere via docker repositories.
4:30 PM					
5:00 PM	Hackerspace	Singapore	Meetup and Hacking		