

Time	Given Name	Family Name	Type	Session Topic	Abstract of talk or project
<b>Saturday March 14</b>					
9:00 AM	Brian	Proffitt	User Meetup/Workshop - 3 hours	oVirt Workshop	We would like to host a full- or half-day oVirt user/admin meetup, colocated with the FOSSAsia event. We would provide speakers to fill the available time, and would just need a room with standard AV facilities.
11:50 AM	Group	Photo	Photo	Group Photo	Group Photo at Blk71
<b>12:00 PM</b>					
1:00 PM	Leap	Sok	Talk of 25 minutes	IT Automation with Ansible	Software configuration management tools are gaining ground everywhere. We've all probably at least heard of puppet, chef or salt by now, however there's a new rising star: Ansible. In this talk we'll talk about the way Ansible approaches configuration management, software deployment, remote execution and other general IT tasks.
1:30 PM	Gerald	Goh	Talk of 25 minutes	Docker + Redhat + Openstack and the glue to hold them together	At Apvera we are focused on delivering services that make use of open source technologies to enhance how end users consume business applications. We primarily use Redhat + Docker + Openstack to enable enterprises to help consume application provisioning with simplicity and portability.
2:00 PM	Masahiro	Nakagawa	Talk of 25 minutes	Fluentd: Unified Logging Infrastructure	Fluentd is an open source data collector to unify logging infrastructure. Used at companies like Nintendo and Slideshare and integrated into projects like Kubernetes (Google's Docker orchestration framework), it is quickly becoming a popular tool to collect, process and manage log streams. In this talk, I will highlight the key features and design principles of Fluentd and explain why many programmers and ops engineers find it useful. Time permitting, I will share the project's roadmap and new features. In 2014, Fluentd maintainers have given talks at RedDotRubyConf, RubyConf, MongoDBWorld, OSCON and RubyKaigi. This year, we are expanding our reach to open source communities outside of Ruby (Fluentd is primarily written in Ruby).
2:30 PM	Brett	Gray	Talk of 25 minutes	MCollective: a power tool for ad hoc management	MCollective is a powerful tool that can be used for ad hoc management in geo-diverse environments. This talk will provide some insight into the tool, including architecture, security and extending.
3:00 PM	Stephan H.	Wissel	Talk of 25 minutes	Run your cloud on Open Standards - and have your cake too	You're a startup, so you live in the cloud. You don't have to abandon Open Standards there and become dependent on a propriety stack to grow your company. The session will guide through OpenStack and CloudFoundry, where, how and when you want to use it and who of the big boys is doing what. Some slides, some demos - lots of questions
3:30 PM	Choon Ming	Goh	Talk of 25 minutes	Deploying OpenNebula, the Open Source Solution for Enterprise Cloud Solution	As enterprises expand their datacenters or internal infrastructure, it is very common for one to use commercial software such as VMWare or OSS solution such as Xen, KVM or OpenStack. However, the reality is that enterprises chose commercial products in the end because of the features made available to them. An open source cloud/infrastructure solution such as OpenNebula provides the features that can rivals VMWare whilst integrating with various other open source technologies will balance up the scale. The talk will show you how OpenNebula works and what are the features available and how it integrates with other Open Source solution in order to rival VMWare or other commercial products.
4:00 PM	Shivji Kumar	Jha	Talk of 25 minutes	Enhanced High Availability using MySQL Group Replication	MySQL Replication provides a solution for High Availability and Read Scale-Out. Replication ensures that data written on one MySQL server is made available on other MySQL servers at runtime in a fast, consistent and fault tolerant manner with minimal impact to the overall performance of the server. Traditionally, MySQL Replication supports a single master and many slaves, and it is either asynchronous or semi-synchronous. Recently, a preview of a new replication plugin for MySQL was released and this is named MySQL Group Replication. This plugin provides multi-master update everywhere capabilities, making it possible to update data, concurrently, on any server in a group. While not a fully synchronous replication solution such as MySQL Cluster, it does provide additional synchronization related to message exchanged between servers in a group. This talk explains how MySQL Group Replication facilitates improves High Availability and simplifies replication and application management.
<b>4:30 PM End</b>					
<b>Sunday March 15</b>					
9:30 AM	Chris	Kruppa	Talk of 25 minutes	How a lean and agile mindset can lead to more innovative products	Most innovative projects fail either because of wrong and not verified assumptions or because of fearing the risk to invest into insecure projects. The lean and agile mindset deals with those fears and offer certain practices to both develop an innovative product and to find fast the market fit. Therefore investments can be reduced and the chance for success increases by factors. In this workshop we will learn some basics of the lean/agile mindset and apply some of the practices introduced by it's leaders. The workshop will be very interactive: It is planned, that the participants will choose the practices they want to learn during this workshop.
10:00 AM	Kuan Yen	Heng	Talk of 25 minutes	Docker in Practice	Docker is a containerization platform that has won over legions of developers who fancy themselves as sysadmins. As the early adopters come to grips with the technology, a set of patterns and best practices start to emerge. Here are some of the lessons we've learnt in optimizing Docker development and deployment workflows for continuous delivery.
10:30 AM	Amita	Sharma	Talk of 25 minutes	Testing - Best Practices in Open Souce World	1. How to contribute in improving quality of open source software 2. How to create test harness for open source software. 3. How to alter your software development model to produce bug-free code. 4. Altering your development model to match SW quality standards. 5. Study the technology, know the feature well and use your experience 6. Know the point when it is time to report a bug 7. How to troubleshoot that buggy situation.

Time	Given Name	Family Name	Type	Session Topic	Abstract of talk or project
11:00 AM	Maor	Lipchuk	Talk of 25 minutes	Disaster Recovery in oVirt	Even the best system administrator cannot always avoid any and every disaster that may plague his data center, but he should have a contingency plan to recover from one - and an administrator that manages his virtual data centers with oVirt is of course no different. This session will cover the new features introduced in oVirt 3.5.0 to handle such scenarios and will showcase how stringing together a set of building blocks can produce a well rounded solution for disaster scenarios.
11:30 AM	Satoshi	Nagayasu	Talk of 25 minutes	PostgreSQL 9.4 and Beyond	Recently, PostgreSQL has succeeded in getting popularity among web developers. One of the reasons is that PostgreSQL is growing not only as a traditional RDBMS, but also as an application development platform with fully supporting ACID properties. In this talk, I would like to introduce new features in the latest release 9.4 (JSON, Analytics, Logical Decoding, etc.) to take advantages of the latest PostgreSQL advancement, and share a quick look in the upcoming release 9.5.
<b>12:00 PM</b>	Lunch			Break	
1:00 PM	Piyush	Verma	Talk of 25 minutes	analysing your CDN usage	CDNylis syncs Amazon Cloudfront logs, which are stored as compressed files in S3 bucket and streams them to InfluxDB, a time series Database with a powerful Querying API. CDNylis is designed to work on millions and millions of records and as-it-happens updates of the CDN usage. You use this for: Understanding how the Bandwidth is being used. Finding out the most popular and most downloadable content. Generate trends for your most popular Videos, Audios, Slides etc. Understand geographical behaviour of the Requests. Amount of Bytes transferred to & fro the Cloudfront distributions. Find out the most profitable referrer from where your content is being accessed. etc.
1:30 PM	Ryusuke	Kajiyama	Talk of 25 minutes	MySQL Cluster as Transactional NoSQL	There's a lot of excitement around NoSQL Data Stores with the promise of simple access patterns, flexible schemas, scalability and High Availability. The downside comes in the form of losing ACID transactions, consistency, flexible queries and data integrity checks. What if you could have the best of both worlds? This session shows how MySQL Cluster provides simultaneous SQL and native NoSQL access to your data - whether a simple key-value API (Memcached), REST, JavaScript, Java or C++. You will hear how the MySQL Cluster architecture delivers in-memory real-time performance, 99.999% availability, on-line maintenance and linear, horizontal scalability through transparent auto-sharding.
2:00 PM	Sameer	Kumar	Talk of 25 minutes	10 Features your Developers are missing when 'Stuck' with proprietary database!	Database is core of any enterprise project. Be it the solution which a bank relies on for daily transactions or be the software which helps your mobile operator track your usage and send you a bill. Every application needs a data store and that is why every developer sooner or later has to be exposed to databases. I have seen mostly when I talk to developers about a database or relational database there are a few or rather just a name that would pop-up in their head. Proprietary databases have become so much a synonym to the term "RDBMS" that developer today hardly know about any other database. In this talk I will discuss about 10 such great features which your developers are missing if they are still stuck with Proprietary Database. I will be discussing 10 awesome features of PostgreSQL- Worlds most Advanced Open Source Database.
2:30 PM	Mayank	Prasad	Talk of 25 minutes	MySQL Performance Schema - Monitoring MySQL Server with Performance Schema and enhancement in MySQL 5.7	The MySQL Performance Schema is a feature for monitoring Oracle's MySQL Server execution at a low level. This versatile and tightly integrated component collects performance and session data from various subsystems within the server during runtime with minimal impact on overall server performance. This feature was introduced in MySQL 5.5 and MySQL 5.6 was released with great enhancements in it for performance monitoring and tuning. In 5.7 milestone releases as well, there are multiple new functionalities have been added. This talk aims to give an overview of Performance Schema and how/what statistical information user is provided with using Performance Schema and how that information can be used for monitoring and trouble shooting to track an issue back to the relevant file and line of code in source file. This talk also explains how Performance Schema gives consolidated view of different server activities behind the scene with the multiple summary tables and how this can be tuned. This talk also aims to give a glimpse of new features/monitoring information added in MySQL 5.7 Milestone releases. This talk also covers the manual configuration of Performance Schema to suite to individual requirement in a specific scenario.
3:00 PM	Aditya	Patawari	Talk of 25 minutes	Running Project Atomic and Docker on Fedora	This session will focus on getting started with Docker with an introduction to Project Atomic. We'll discuss why Docker can be a better choice than LXC and virtual machines in many scenarios. We'll discuss rpm-ostree and its advantages followed by running a Fedora Atomic host feeding it cloud-init data. We'll also create custom docker images from Fedora image which can be shipped anywhere via docker repositories. This talk is targeted to many database users which includes DBA, Database Application developers, Database developers, Students. This talk is to give overview of MySQL Performance Schema and also how/what information can be collected using it and how this information could be used further for debugging/trouble shooting Database users' issue. Takeaways from the talk: - Overview of tuning and monitoring MySQL Server activities using Performance Schema. - Understanding of how to tailor the Performance Schema to track the data you want. - Demo with a use case to show how to diagnosing problems with the Performance Schema. - What's new in Performance Schema in MySQL 5.7 (milestone) releases.
3:30 PM	Lyle	Kozloff	Talk of 25 minutes	Configuration Management Systems on the Desktop	
4:00 PM	Anisha	Narang	Talk of 25 minutes	All about web application testing	Web application testing is all about validating the user flow and checking the functionality available to the end user. We would explore some tools we can use to test the front end of a web application. Talk a little about tools like Selenium, Wait and see if we can use these along with the programming language of our choice. Also, I would be considering some of the common issues faced while front end testing comes to picture and discuss some best practices.

Time	Given Name	Family Name	Type	Session Topic	Abstract of talk or project
4:25 PM	Hong Phuc / Mario	Dang / Behling	Talk	Wrap up	See you 2016!
<b>4:35 PM</b>	<b>End</b>				
<b>5:00 PM</b>	<b>Hackerspace</b>	<b>Singapore</b>	<b>Meetup and Hacking</b>		Address: Hackerspace, 344B King George's Avenue, Singapore 208576
<b>Venue</b>	<b>Singtel Innov8</b>		Address	71 Ayer Rajah Crescent #02-22, Singapore 139951, Phone: +65 6838 4686, Fax: +65 6872 8456	
Room	Seminar Room		How to get here	Alight at One North MRT. Go up on the escalator to Ground floor level and exit at Fusionopolis. Walk over the street towards Blk71 and enter the building. On the first floor you find Plug-In at #02-22.	
Link	<a href="http://innov8.singtel.com">http://innov8.singtel.com</a>				
Map	<a href="https://www.google.com.sg/maps/place/SingTel+Innov8+Pte+Ltd/@1.29677,103.786914,17z/data=!3m1!1e4m2!3m1!1s0x31da1a4fd6113e5f:0x3c1e4af504244a15?">https://www.google.com.sg/maps/place/SingTel+Innov8+Pte+Ltd/@1.29677,103.786914,17z/data=!3m1!1e4m2!3m1!1s0x31da1a4fd6113e5f:0x3c1e4af504244a15?</a>				